

INTRODUCTION:

Aging, an inevitable and intricate biological process, has always intrigued humanity. Understanding age not only helps us gain insight into the human condition but also has profound implications across various fields, such as healthcare, entertainment, marketing, and security. Deep learning, a subset of machine learning, has emerged as a potent tool to unravel age-related mysteries by analyzing patterns, textures, and features in images.

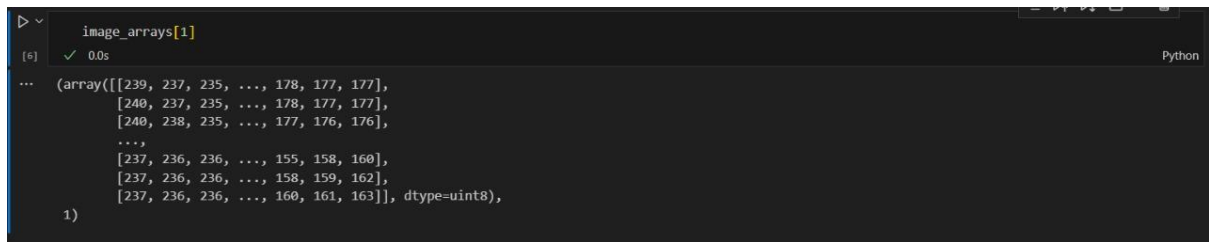
Our project embarks on a journey to explore the capabilities of deep learning in age estimation, seeking to answer fundamental questions such as:

1. How accurately can deep learning algorithms predict a person's age based on facial features and images?
2. What role do convolutional neural networks (CNNs), recurrent neural networks (RNNs), and other deep learning architectures play in age detection?
3. What datasets are available for training and testing age detection models, and how do they contribute to the accuracy of predictions?

DATA SOURCE:

<https://drive.google.com/drive/folders/1E9m9dZYLga9kc9NGPHfZa75JNJgpgv3M>

The folders in the dataset are named according to the age of people in it.



```
image_arrays[1]
Out[6]: array([[239, 237, 235, ..., 178, 177, 177],
        [240, 237, 235, ..., 178, 177, 177],
        [240, 238, 235, ..., 177, 176, 176],
        ...,
        [237, 236, 236, ..., 155, 158, 160],
        [237, 236, 236, ..., 158, 159, 162],
        [237, 236, 236, ..., 160, 161, 163]], dtype=uint8),
1)
```

Age Grouping:

We created age groups, such as 0-5, 6-12, 13-18, 19-30, 31-45, 46-60, 61-80, and 81-95, in our machine learning model for simplifying the modeling process, enhancing interpretability, and aligning with the practical significance of these groups in our application. These age groups were chosen after thorough experimentation and analysis to maximize model accuracy and performance while ensuring compliance with ethical considerations. They facilitate decision-making, scalability, and maintenance, making them a suitable approach for addressing our project's specific needs and objectives.

DATA AUGMENTATION:

- Splitting Data:



```
from sklearn.model_selection import train_test_split

train_data, temp_data = train_test_split(image_arrays, test_size=0.3, random_state=43)

val_data, test_data = train_test_split(temp_data, test_size=0.5, random_state=42)
```

- Separating Dependent and independent variable:

```
Separating Dependent and Independent variable
markdown

x_train = [item[0] for item in train_data]
y_train = [item[1] for item in train_data]

x_test = [item[0] for item in test_data]
y_test = [item[1] for item in test_data]

x_val = [item[0] for item in val_data]
y_val = [item[1] for item in val_data]

[8] ✓ 0.0s Python
```

```
Looking structure of data
Python

Loading...
x_train[0]
[9] ✓ 0.0s Python

... array([[ 76,  74,  68, ...,  77,  82,  85],
          [ 77,  73,  68, ...,  76,  80,  83],
          [ 77,  73,  68, ...,  75,  77,  79],
          ...,
          [ 55,  53,  50, ...,  48,  72, 100],
          [ 54,  56,  55, ...,  46,  68,  95],
          [ 55,  58,  60, ...,  41,  62,  89]], dtype=uint8)

y_train[0]
[10] ✓ 0.0s Python

... 5

x_test[0]
[11] ✓ 0.0s Python

... array([[ 2,  1,  1, ..., 30, 32, 33],
          [ 2,  1,  1, ..., 29, 31, 32],
          [ 1,  1,  1, ..., 28, 29, 30],
          ...,
          [23, 21, 19, ...,  8, 13, 16],
          [23, 21, 19, ...,  8, 13, 16],
          [22, 21, 19, ...,  8, 13, 15]], dtype=uint8)
```

```
len(x_train)
[13] ✓ 0.0s Python

... 6828

len(x_test)
[14] ✓ 0.0s Python

... 1464

len(x_val)
[15] ✓ 0.0s Python

... 1463
```

- Resizing image:

```
x_train[0].shape
[18] ✓ 0.0s Python

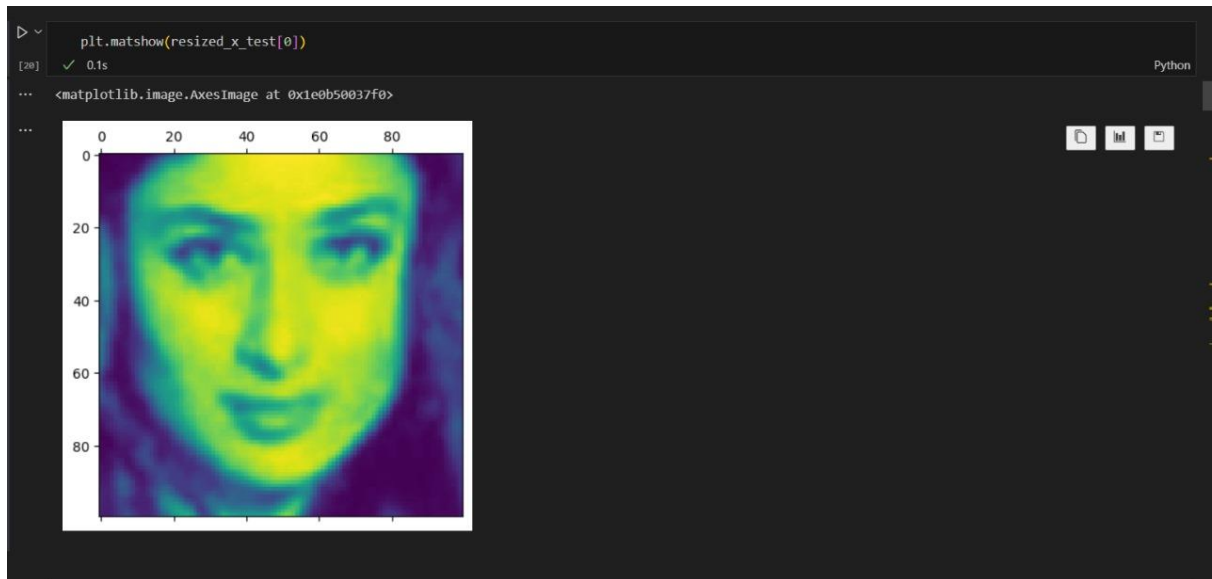
... (200, 200)
```

```
Resizing images pixel  
[ ] Python  
new_size = (100, 100)  
resized_x_train = []  
  
for image in x_train:  
    resized_image = cv2.resize(image, new_size)  
    resized_x_train.append(resized_image)  
  
resized_x_train = np.array(resized_x_train)  
  
print(resized_x_train.shape)  
  
resized_x_test = []  
  
for image in x_test:  
    resized_image = cv2.resize(image, new_size)  
    resized_x_test.append(resized_image)  
  
resized_x_test = np.array(resized_x_test)  
  
resized_x_val = []  
  
for image in x_val:  
    resized_image = cv2.resize(image, new_size)  
    resized_x_val.append(resized_image)  
  
resized_x_val = np.array(resized_x_val)  
[17] ✓ 0.2s Python  
... (6828, 100, 100)  
Ln 1, Col 10 Spaces: 4 CRLF Cell 16 of 71
```

- Normalizing Data Set

```
Converting to numpy array  
markdown  
x_train_np = np.array(x_train)  
x_test_np = np.array(x_test)  
x_val_np = np.array(x_val)  
y_train_np = np.array(y_train)  
y_test_np = np.array(y_test)  
y_val_np = np.array(y_val)  
Loading...  
x_train_normalized = resized_x_train / 255.0  
x_test_normalized = resized_x_test / 255.0  
x_val_normalized = resized_x_val / 255.0  
[32] ✓ 0.5s Python
```

- Visualizing Dataset:



```
x_test_normalized[0]
```

[23] ✓ 0.0s

... array([[0.00784314, 0.00392157, 0.00784314, ..., 0.08627451, 0.11372549,
0.1254902],
[0.00392157, 0.00392157, 0.00392157, ..., 0.0745098 , 0.10196078,
0.11372549],
[0. , 0. , 0.00392157, ..., 0.07058824, 0.09411765,
0.10980392],
...,
[0.09019608, 0.07843137, 0.07843137, ..., 0.02745098, 0.03529412,
0.06666667],
[0.08627451, 0.0745098 , 0.0745098 , ..., 0.02352941, 0.03137255,
0.05882353],
[0.08627451, 0.0745098 , 0.0745098 , ..., 0.02352941, 0.03137255,
0.05490196]])

```
x_test_normalized.shape
```

[24] ✓ 0.0s

... (1464, 100, 100)

```
x_train_normalized.shape
```

[25] ✓ 0.0s

... (6828, 100, 100)

- Flattering Dataset:

```
Flattering_X_dataset
```

markdown

```
X_train_flattened = x_train_normalized.reshape(len(x_train_normalized), 100*100)  
X_test_flattened = x_test_normalized.reshape(len(x_test_normalized), 100*100)  
X_val_flattened = x_val_normalized.reshape(len(x_val_normalized), 100*100)
```

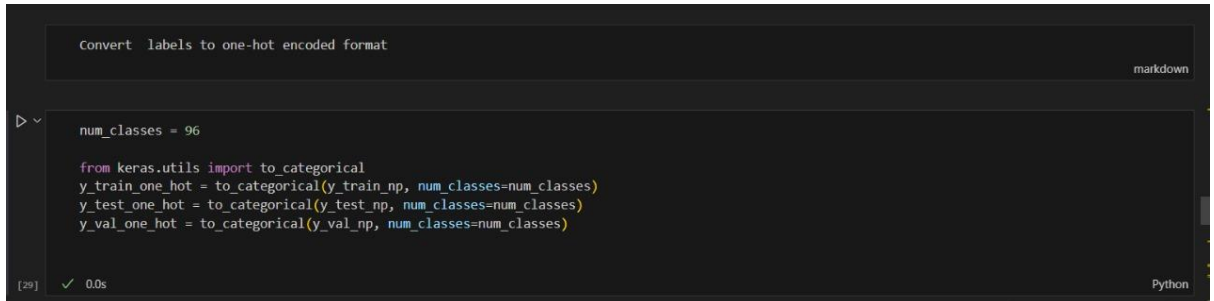
[26] ✓ 0.0s

```
x_train_flattened.shape
```

[27] ✓ 0.0s

... (6828, 10000)

- Convert labels to one-hot encoded format:



Convert labels to one-hot encoded format

```
num_classes = 96

from keras.utils import to_categorical
y_train_one_hot = to_categorical(y_train_np, num_classes=num_classes)
y_test_one_hot = to_categorical(y_test_np, num_classes=num_classes)
y_val_one_hot = to_categorical(y_val_np, num_classes=num_classes)
```

[29] ✓ 0.0s Python

- Reshaping Dataset:



Reshaping x_dataset

```
x_train_resaped = x_train_flattened.reshape(-1, 100, 100, 1)
x_val_resaped = x_val_flattened.reshape(-1, 100, 100, 1)
x_test_resaped = x_test_flattened.reshape(-1, 100, 100, 1)
```

[30] ✓ 0.0s Python

MODEL ARCHITECTURE:

```
CNN Architecture

import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense
from tensorflow.keras.regularizers import l2

model = Sequential([
    Conv2D(32, kernel_size=(3, 3), input_shape=(100, 100, 1), activation='relu'),
    MaxPooling2D(pool_size=(2, 2)),
    Conv2D(256, kernel_size=(3, 3), activation='relu'),
    MaxPooling2D(pool_size=(2, 2)),
    Conv2D(512, kernel_size=(3, 3), activation='relu'),
    MaxPooling2D(pool_size=(2, 2)),
    Conv2D(256, kernel_size=(3, 3), activation='relu'),
    MaxPooling2D(pool_size=(2, 2)),
    Flatten(),
    Dense(128, activation='relu'),
    Dense(9, activation='softmax')
])
optimizer = Adam(learning_rate=0.001)
model.compile(optimizer=optimizer,
              loss='categorical_crossentropy',
              metrics=['accuracy'])
```

[35] ✓ 0.9s Python

MODEL TRAINING:

```
Model training | ✓ 📄 ... 🗑️
Code Markdown markdown

from tensorflow.keras.callbacks import EarlyStopping

early_stopping = EarlyStopping(
    monitor='val accuracy',
    patience=10,
    restore_best_weights=True
)

history = model.fit(
    x_train_resaped,
    y_train_one_hot,
    epochs=100, batch_size=16,
    validation_data=(x_val_resaped, y_val_one_hot),
    callbacks=[early_stopping]
)
```

[35] ✓ 3m 5.0s Python

```
1 Epoch 1/100
2 427/427 [=====] - 15s 26ms/step - loss: 1.6532 - accuracy: 0.3650 - val_loss: 1.3653 - val_accuracy: 0.4675
3 Epoch 2/100
4 427/427 [=====] - 10s 23ms/step - loss: 1.2206 - accuracy: 0.5267 - val_loss: 1.1111 - val_accuracy: 0.5379
5 Epoch 3/100
6 427/427 [=====] - 10s 23ms/step - loss: 1.0810 - accuracy: 0.5709 - val_loss: 1.0469 - val_accuracy: 0.5830
7 Epoch 4/100
8 427/427 [=====] - 10s 23ms/step - loss: 0.9742 - accuracy: 0.6112 - val_loss: 1.0894 - val_accuracy: 0.5783
9 Epoch 5/100
10 427/427 [=====] - 10s 23ms/step - loss: 0.9039 - accuracy: 0.6353 - val_loss: 1.0164 - val_accuracy: 0.5762
11 Epoch 6/100
12 427/427 [=====] - 10s 23ms/step - loss: 0.8271 - accuracy: 0.6658 - val_loss: 1.0028 - val_accuracy: 0.6049
13 Epoch 7/100
14 427/427 [=====] - 10s 23ms/step - loss: 0.7633 - accuracy: 0.6861 - val_loss: 1.0330 - val_accuracy: 0.5803
15 Epoch 8/100
16 427/427 [=====] - 10s 23ms/step - loss: 0.7036 - accuracy: 0.7154 - val_loss: 1.0089 - val_accuracy: 0.6186
17 Epoch 9/100
18 427/427 [=====] - 10s 23ms/step - loss: 0.6208 - accuracy: 0.7525 - val_loss: 1.0100 - val_accuracy: 0.6124
19 Epoch 10/100
20 427/427 [=====] - 10s 23ms/step - loss: 0.5513 - accuracy: 0.7743 - val_loss: 1.0989 - val_accuracy: 0.6104
21 Epoch 11/100
22 427/427 [=====] - 10s 23ms/step - loss: 0.4839 - accuracy: 0.8084 - val_loss: 1.1833 - val_accuracy: 0.6042
23 Epoch 12/100
24 427/427 [=====] - 10s 23ms/step - loss: 0.4131 - accuracy: 0.8363 - val_loss: 1.3691 - val_accuracy: 0.5878
25 Epoch 13/100
26 427/427 [=====] - 10s 23ms/step - loss: 0.3429 - accuracy: 0.8626 - val_loss: 1.3977 - val_accuracy: 0.5933
27 Epoch 14/100
28 427/427 [=====] - 10s 23ms/step - loss: 0.2900 - accuracy: 0.8903 - val_loss: 1.6242 - val_accuracy: 0.6124
29 Epoch 15/100
30 427/427 [=====] - 10s 23ms/step - loss: 0.2614 - accuracy: 0.9001 - val_loss: 1.5280 - val_accuracy: 0.5871
31 Epoch 16/100
32 427/427 [=====] - 10s 23ms/step - loss: 0.2249 - accuracy: 0.9145 - val_loss: 1.8085 - val_accuracy: 0.5892
33 Epoch 17/100
34 427/427 [=====] - 10s 23ms/step - loss: 0.2104 - accuracy: 0.9209 - val_loss: 2.1422 - val_accuracy: 0.5954
35 Epoch 18/100
36 427/427 [=====] - 10s 23ms/step - loss: 0.1764 - accuracy: 0.9348 - val_loss: 2.0044 - val_accuracy: 0.5960
37
```

LOSS AND ACCURACY CURVE:

```
Plotting Loss and Accuracy curves
markdown

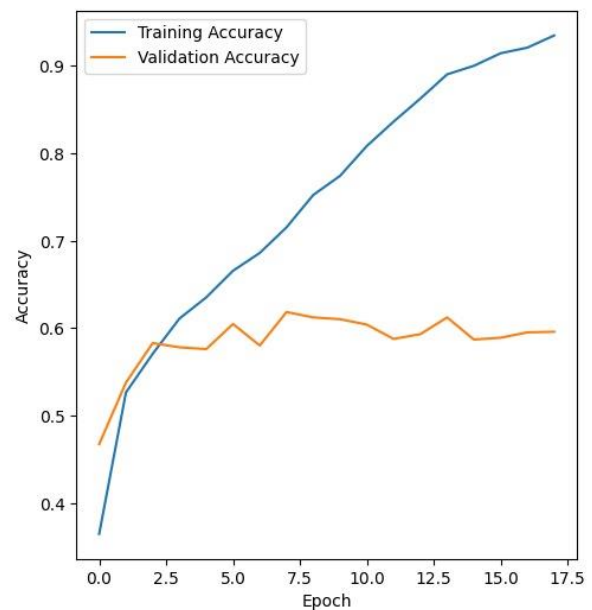
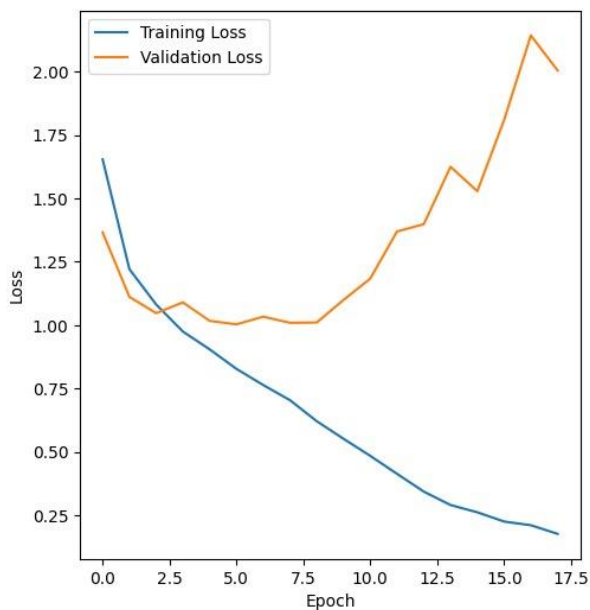
train_loss = history.history['loss']
val_loss = history.history['val_loss']
train_acc = history.history['accuracy']
val_acc = history.history['val_accuracy']

plt.figure(figsize=(12, 6))
plt.subplot(1, 2, 1)
plt.plot(train_loss, label='Training Loss')
plt.plot(val_loss, label='Validation Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(train_acc, label='Training Accuracy')
plt.plot(val_acc, label='Validation Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()

plt.show()
```

[41] ✓ 02s Python



EVALUATION:

```
Evaluation on test data  
[42] ✓ 0.0s  
x_test_resaped = x_test_resaped.reshape(-1, 100, 100, 1)  
[43] ✓ 1.1s  
y_predict = model.predict(x_test_resaped, batch_size=16)  
... 92/92 [=====] - 1s 9ms/step  
+ Code + Markdown  
[44] ✓ 1.1s  
model.evaluate(x_test_resaped, y_test_one_hot, batch_size=16)  
... 92/92 [=====] - 1s 10ms/step - loss: 1.0774 - accuracy: 0.6093  
... [1.077383279800415, 0.6092896461486816]
```

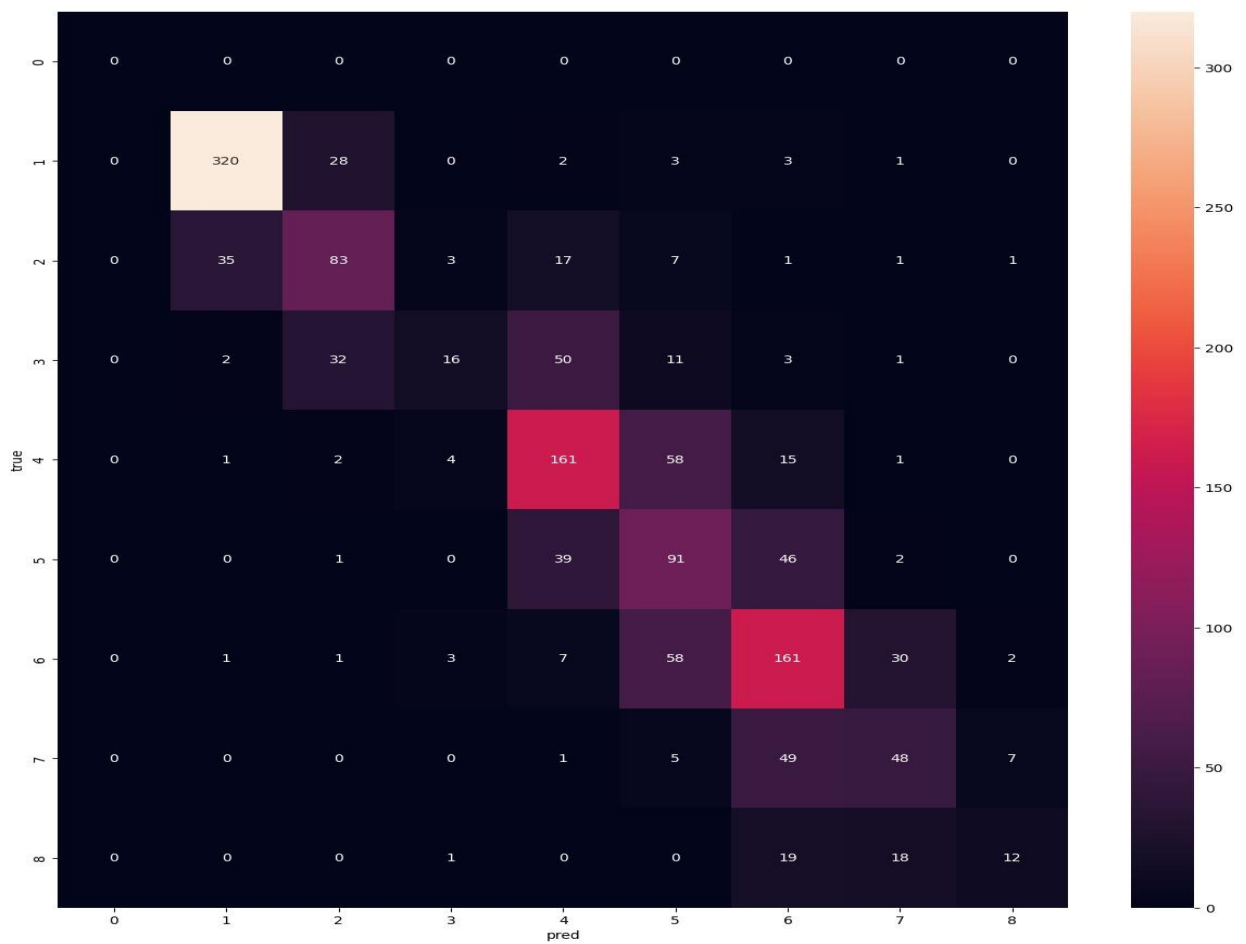
CONFUSION MATRIX:

```
Plotting confusion Matrix

y_predict_labels=[np.argmax(i) for i in y_predict]
cn=tf.math.confusion_matrix(labels=y_test_np,predictions=y_predict_labels)
[45] ✓ 0.0s Python

import seaborn as sn
plt.figure(figsize=(15,15))
sn.heatmap(cn, annot=True, fmt='d')
plt.xlabel('pred')
plt.ylabel('true')
[46] ✓ 0.4s Python

... Text(158.22222222222223, 0.5, 'true')
```



SAVING THE MODEL:

```
Saving the model

def predict_with_model(model, input_data):
    return model(input_data)

# Save the model
model.save("saved_model")

[ ] Python

import tensorflow as tf

loaded_model = tf.keras.models.load_model("saved_model")

[ ] Python
```

Real Time Testing

```
Real life Testing

image = cv2.imread('img_name.jpg')

img_gray = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)

resized_image = cv2.resize(img_gray, (100,100))
image_np = np.array(resized_image)

resized_image = np.reshape(image_np, (-1, 100, 100, 1))

img_normalized = resized_image/ 255.0

prediction=loaded_model.predict(img_normalized)
max_position = np.argmax(prediction)

print("YOUR AGE IS: ")

if(max_position==1):
    print("0-5")
elif(max_position==2):
    print("6-12")
elif(max_position==3):
    print("13-18")
elif(max_position==4):
    print("19-30")
elif(max_position==5):
    print("31-45")
elif(max_position==6):
    print("46-65")
elif(max_position==7):
    print("66-80")
else:
    print(">81")
```